

Siddhant Agarwal

[Google Scholar](#) | siddhant@cs.utexas.edu | [LinkedIn](#) | agarwalsiddhant10.github.io

EDUCATION

The University of Austin at Texas

Ph.D. in Computer Science

Advisors: Professor Amy Zhang and Professor Peter Stone

Austin, USA

August 2022 - May 2027 (expected)

Indian Institute of Technology Kharagpur

B.Tech. (Hons.) + M.Tech. (Dual Degree) in Computer Science and Engineering

Department Rank 1 (GPA 9.90/10.00)

Kharagpur, India

July 2017 - June 2022

FIELDS OF INTEREST

Unsupervised RL, GCRL, Representation Learning, Robotics, LLM Finetuning & Reasoning

PUBLICATIONS

(* denotes equal contribution, [†] denotes equal advising)

A Unifying Perspective on Unsupervised Reinforcement Learning Algorithms

Siddhant Agarwal*, Caleb Chuck*, Harshit Sikchi*, Jiaheng Hu, Max Rudolph, Scott Niekum[†], Peter Stone[†], Amy Zhang[†]

Submitted *14th International Conference on Learning Representations, 2026*

RLC Workshop on Reinforcement Learning Beyond Rewards, 2025

Regularized Latent Dynamics Prediction is a Strong Baseline For Behavioral Foundation Models

Pranaya Jajoo, Harshit Sikchi*, **Siddhant Agarwal***, Amy Zhang, Scott Niekum, Martha White

Submitted *14th International Conference on Learning Representations, 2026*

RLC Workshop on Reinforcement Learning Beyond Rewards, 2025

RL Zero: Direct Policy Inference from Language Without In-Domain Supervision

Harshit Sikchi*, **Siddhant Agarwal***, Pranaya Jajoo*, Samyak Parajuli*, Caleb Chuck*, Max Rudolph*, Peter Stone[†], Amy Zhang[†], Scott Niekum[†]

39th International Conference on Neural Information Processing Systems(NeurIPS), 2025

Proto Successor Measure: Representing the space of all possible solutions of Reinforcement Learning

Siddhant Agarwal*, Harshit Sikchi*, Peter Stone[†], Amy Zhang[†]

42th International Conference on Machine Learning (ICML), 2025

Reinforcement Learning Within the Classical Robotics Stack: A Case Study in Robot Soccer

Adam Labiosa*, Zhihan Wang*, **Siddhant Agarwal**, William Cong, Geethika Hemkumar, Abhinav Narayan Harish, Benjamin Hong, Josh Kelle, Chen Li, Yuhao Li, Zisen Shao, Peter Stone, Josiah P. Hanna
IEEE Conference on Robotics and Automation (ICRA), 2025

Robot Air Hockey: A Manipulation Testbed for Robot Learning with Reinforcement Learning

Caleb Chuck*, Carl Qi*, Michael J. Munje*, Shuzhe Li*, Max Rudolph*, Chang Shi*, **Siddhant Agarwal***, Harshit Sikchi*, Abhinav Peri, Sarthak Dayal, Evan Kuo, Kavan Mehta, Anthony Wang, Peter Stone, Amy Zhang, Scott Niekum,
ICRA Workshop Manipulation Skills, 2024

f-Policy Gradients: A General Framework for Goal Conditioned RL using *f*-Divergences

Siddhant Agarwal, Ishan Durugkar, Peter Stone, Amy Zhang

37th International Conference on Neural Information Processing Systems(NeurIPS), 2023

Behavior Predictive Representations for Generalization in Reinforcement Learning

Siddhant Agarwal, Aaron Courville, Rishabh Agarwal

The Multi-Disciplinary Conference on Reinforcement Learning and Decision Making (RLDM), 2022,

NeurIPS 2021 EcoRL and DeepRL workshops

Reinforcement Explanation Learning

Siddhant Agarwal, Owais Iqbal, Sree Aditya Buridi, Madda Manjusha, Abir Das

NeurIPS 2021 workshop on eXplainable AI approaches for debugging and diagnosis [Oral]

Poisoned Classifiers are not only backdoored, they are fundamentally broken

Mingie Sun, **Siddhant Agarwal**, Zico Kolter

ICLR 2021 workshop on Security and Safety in Machine Learning Systems

Learning to Deceive Knowledge Graph Augmented Models via Targeted Perturbations

Mrigank Raman, Aaron Chan*, **Siddhant Agarwal***, Peifeng Wang, Hansen Wang, Sungchul Kim, Ryan Rossi, Handong Zhao, Nedim Lipka, Xiang Ren

International Conference on Learning Representations (ICLR), 2021

Traffic Sign Classification using Hybrid HOG-SURF Features using Convolutional Neural Networks

Rishabh Madan*, Deepank Agrawal*, Shreyas Kowshik*, Harsh Maheshwari*, **Siddhant Agarwal***, Debasish Chakraborty (* denotes equal contribution-dice roll)

International Conference on Pattern Recognition Applications and Methods, Prague 2019

EXPERIENCE

Applied Research Scientist Intern

Amazon

May 2024-April 2025

New York

Developed a principled exploration strategy to collect online samples to provably improve offline RL datasets by reducing the sub-optimality gap in the offline dataset. Integrated this algorithm to improve the Supply Chain buying policy.

Data Engineering Intern

Goldman Sachs

June 2021 - July 2021

India

Developed a python client to use the services of an ETL platform to quickly transfer data from pandas dataframe to a MemSQL store. Created RESTful server endpoints for the ETL platform and accessed them from the client.

Research Intern

LocusLab, Carnegie Mellon University, Advisor: Zico Kolter

May 2020 - October 2020

Pittsburg, USA

Showed that *backdoored classifiers* are fundamentally broken and can be attacked by anyone rather than only the adversary.

Research Intern

INKLab, University of Southern California, Advisor: Xiang Ren

March 2020 - May 2020

Los Angeles, USA

Trained DQN policies to heavily perturb KGs without changing the downstream performance of the different KG models. Was able to deceive MHGRN, GRN and RN (on CSQA and OBQA) and KGNN and RippleNet (on MovieLens and Music).

Autonomous Ground Vehicles Research Group

Indian Institute of Technology Kharagpur

March 2018 - March 2022

Kharagpur, India

Worked to develop a category 4 autonomous ground vehicle. Developed traffic sign classification system for Indian traffic signs, lane segmentation systems, planning using Frenet Serret Local Planner and Time Elastic Bands and Localization systems using Extended Kalman Filters.

INVITED TALKS

The Causal Reinforcement Learning Workshop, RL Conference

Learning Generalist Agents using Causal MDP Abstractions

2025

The International Summer School on Situational Awareness in Cognitive Technologies

with Prof. Abir Das, hands-on session with on explainable AI

2021

COMPETITIONS

RoboCup Standard Platform League *First Place in Challenge Shield, 2024*

2023, 2024

Prepared a team of 5 Nao robots to compete in a soccer tournament which involved developing the basic motion system like walking and kicking along with localization to estimate the agent's position in the field. Used RL to train behavior strategies in the game to score goals.

Bosch Traffic Sign Detection, Inter IIT Tech Meet *Bronze*

2021

Developed a web-based tool to add real-life data augmentations on a traffic sign dataset, train classifiers, and debug them. The tool helps the users to visualize the mistakes of the model and debug the model.

Intelligent Ground Vehicles Competition, Oakland University, Michigan, *Runners Up*

2018 & 2019

Developed an autonomous ground vehicle for a constraint environment using simple Computer Vision and Image Processing techniques for Lane navigation and Obstacle avoidance, TEB local planner for planning. First team to qualify.

TEACHING EXPERIENCES AND ORGANIZING SERVICES

Graduate Teaching Assistant

* CS394R Reinforcement Learning, The University of Texas at Austin	2024
* CS343 Artificial Intelligence, The University of Texas at Austin	2022
* CS60010 Deep Learning, IIT Kharagpur	2022
* CS60077 Reinforcement Learning, IIT Kharagpur	2021

Workshop on Reinforcement Learning Beyond Rewards: Ingredients for Developing Generalist Agents

at The RL Conference, Edmonton, *Organizer*

Member of the core organization team for the workshop, worked on inviting speakers, schedule, website, accepted paper list, spotlights and overall management of the workshop.

Workshop on Reinforcement Learning Beyond Rewards at The RL Conference, Amherst, *Organizer*

Member of the core organization team for the workshop, helped to ideate the workshop, worked on the invited speaker invitations, schedule, website, accepted paper list, spotlights and overall management of the workshop.

Dynamic Neural Networks meet Computer Vision Workshop at CVPR, *Volunteer & Reviewer*

Volunteered to help the organizers with the logistics of the workshop, and created the workshop posters.

Reviewer

ICLR 2026, NeurIPS 2025 (**Top Reviewer**), ICML 2025, ICLR 2025, RLC 2025, NeurIPS 2024, ICLR 2024, NeurIPS 2023
Goal Conditioned Reinforcement Learning workshop, NeurIPS 2023, CVPR 2021 workshop on Dynamic Neural Networks meet Computer Vision

ACHIEVEMENTS

Top Reviewer, NeurIPS

2025

Most Creative Poster at the Texas Robotics Symposium

2023

Graduate Scholarship from ECE Dept, UT Austin

2023

Institute Silver Medal for securing Rank 1 in the Department of Computer Science and Engineering

2022

GATE Scholarship for Teaching Assistantship

2021

IUSSTF - Viterbi Scholarship - *Only ten students selected from India*

2020

Class of 1970 Alumni (US) Association Prize for *best student in order of merit* among third year B.Tech.(Hons.)/B.Arch.(Hons.) courses securing highest CGPA

2019

IIT Kharagpur Alumni (California Chapter) Award for being the *best student in order of merit* among third year B.Tech.(Hons.)/B.Arch.(Hons.) courses securing highest CGPA

2019

Student Par Excellence Award by Computer Science & Engineering Dept

2018

KVPY Fellowship Award - *Among the top 0.6% of the applicants*

2017

OUTREACH

Texas Robot Parade

2023, 2024, 2025

Demonstrated the Naos from our RoboCup Standard Platform League Team that plays soccer in a competitive tournament.

Directed Reading Program, *Mentor*

2023

Organize a reading group on reinforcement learning for undergraduate students to get them involved in RL research